

7

The Disinformation Machine

Alex Stamos, who was at the time head of security at Facebook, wrote on a blog post, “We have found approximately \$100,000 in ad spending from June of 2015 to May of 2017—associated with roughly 3,000 ads—that was connected to about 470 inauthentic accounts and Pages in violation of our policies. Our analysis suggests these accounts and Pages were affiliated with one another and likely operated out of Russia.” The majority of ads did not directly mention either U.S. presidential candidate but “appeared to focus on amplifying divisive social and political messages across the ideological spectrum—touching on topics from LGBT matters to race issues to immigration to gun rights.” Stamos’s statement did not say whether this would be the last or only revelation of Russian interference via Facebook. Nor did he reveal how many Facebook users saw the advertisements. The ads, according to an expert on Facebook’s advertising system who spoke to BuzzFeed reporters, likely were seen by between twenty-three million and seventy million people, based on the \$100,000 ad buy alone.¹

Stamos announced this revelation as a supplement to a report that he had coauthored in April 2017. That report, in the wake of the 2016 U.S. election and the growing controversies over the proliferation of false news stories and other forms of propaganda, had conceded that Facebook had hosted and delivered significant amounts of suspect content that served to undermine deliberation and democratic practice. “Our mission is to give people the power to share and make the world more open and connected,” the Facebook report read. “Yet it is important that we acknowledge and take steps to guard against the risks that can arise in online communities like ours. The reality is that not everyone shares our vision, and some will seek to undermine it—but we are in a position to help constructively shape the emerging information ecosystem by ensuring our platform remains a safe and secure environment for authentic civic engagement.” The report did not quantify or offer any examples of what it called “information operations” by states and nonstate actors. It just said that there had been some. The report also did not specify what Facebook could do to “constructively shape the emerging information ecosystem” beyond pledging to deploy machine learning to recognize accounts that seem automated. “In brief, we have had to expand our security focus from traditional abusive behavior,

In September 2016 Facebook revealed that advertising accounts based in Russia had precisely targeted advertisement at segments of American voters with propaganda intended to undermine support for Hillary Clinton’s presidential campaign. Immediately speculation renewed about the depth of Russian meddling in American democracy. Was Cambridge Analytica involved with Russia? Who supplied the data that these ad accounts might have used to target so many Americans? Was there a connection between these ads and the Trump campaign? While much of the discussion and speculation focused on Russia, it became immediately clear from that moment that Facebook has contributed to—and profited from—the erosion of democratic practice and norms in the United States and elsewhere. The audacity of a hostile foreign power trying to influence American voters rightly troubles us. But it should trouble us more that Facebook makes such manipulation so easy and renders political ads exempt from the basic accountability and transparency that healthy democracy demands.

such as account hacking, malware and financial scams, to include more subtle and insidious forms of misuse, including attempts to manipulate civic discourse and deceive people.²²

The September revelation, which came as government officials began raising questions about the internal data Facebook had that might show whether the Trump campaign had actively colluded with those within or working for the Russian government, was the first specific Facebook disclosure of “information operations.” The specific offending elements, according to Stamos, were what the advertising industry calls “dark-post ads,” seen only by the narrow intended audience, obscured by the flow of posts within a Facebook News Feed, and ephemeral. Facebook calls its dark-post service Unpublished Page Post Ads. The service is popular among advertisers for its low cost, ease of use, efficiency, effectiveness, and responsiveness. Facebook gives rich and instant feedback to advertisers, allowing them to quickly tailor new ads to improve outcomes or more granularly customize messages. There is nothing mysterious or untoward about the system itself, as long as it’s being used for commerce instead of politics.³

One week after Facebook revealed the presence of these Russian-based advertisers, it admitted that Russian operatives had put up event pages on Facebook to attract attendees to anti-immigrant rallies in the United States. Another report showed that Russians had put up a Facebook page encouraging Texas to secede from the United States. Three months earlier Reuters had reported that a Russian think tank controlled by the Kremlin, the Russian Institute for Strategic Studies, had generated an elaborate plan to influence the 2016 U.S. presidential election. The document recommended the Kremlin launch a propaganda campaign via social media and Russian state-backed global news outlets to encourage U.S. voters to elect a president who would take a softer line toward Russia. As the campaign rolled on and it looked as if Clinton would win easily, the institute proposed that the social media propaganda effort turn toward undermining faith in the American electoral system by spreading false stories of voter fraud. Back in late November 2016 the *Washington Post* had reported on internet experts warning that Russian-based organizations and companies had seeded Facebook with disinformation intended to undermine faith in American

democracy. These earlier reports helped put pressure on Facebook to finally reveal what it knew. As of September 2017, Facebook had revealed only a slim portion of what others have alleged.⁴

As it turns out, we need not spend time wondering who gave these Russian firms the data needed to target specific American audiences. Russian operatives did not require any help from Trump or Cambridge Analytica. Facebook does all the targeting work in-house. Facebook officials, who seem not to grasp the difference between commercial advertising and political propaganda, have repeatedly refused to reveal details of or data from political advertising campaigns. Rob Sherman, deputy chief privacy officer for Facebook, told Reuters in June 2017 that Facebook holds such data in strict confidence because it is “sensitive,” the equivalent of trade secrets. “In many cases, [advertisers will] ask us, as a condition of running ads on Facebook, not to disclose those details about how they’re running campaigns on our service,” Sherman said.⁵

In other words, those who purchase political ads on Facebook enjoy far greater respect for privacy than those of us who use the service every day. Our records of web use, purchases, locations, and interactions with friends and family are all mined for the benefit of advertisers. This pervasive level of surveillance is unmatched in human history. Facebook has created one of the most profitable advertising machines ever created—one that has drained many millions in revenue away from the very news organizations we depend on to foster democratic deliberation and accountability. When a campaign or interest group buys an ad in a newspaper or on a cable television channel the ad could be seen by broad segments of society. It’s also expensive to run even one version of an ad, let alone a dozen different versions. And there is usually a persistent record of the televised advertisement. In the case of federal campaigns in the United States, the 2002 McCain-Feingold Act requires candidates to state they approve of an ad and thus take responsibility for its content. The Federal Election Commission has, so far, neglected to extend the regulatory framework that governs political advertisements on television or radio to web-based or social-media-based political ads.

That tradition of accountability and transparency does not matter to Facebook. Ads on Facebook meant for twenty- to thirty-year-old home-owning Latino men in Virginia would not be viewed by anyone outside that niche. The ads would be ephemeral. Ads could promote falsehoods or misinformation. No one could respond to or even question the claims made in such ads. No one could criticize a group or campaign for its practices or run a response ad. And there would be no public record of campaign themes, arguments, and strategies for present-day political scientists or future historians to analyze. The potential for abuse is vast. An ad could falsely accuse a candidate of the worst malfeasance forty-eight hours before election day and the victim would have no way of knowing it even happened. Ads could stoke ethnic or gender hatred and no one could prepare or respond before serious harm occurs. This should not surprise us. Anyone can deploy Facebook ads. They are affordable and easy. That's one reason that Facebook has grown so quickly into a financial supernova, taking in \$27.6 billion in revenue in 2016 by serving up the attention of two billion Facebook users across the globe.⁶

Daniel Kreiss, a communication scholar at the University of North Carolina, proposes that services such as Facebook, Twitter, and YouTube maintain a repository of campaign ads so that regulators, scholars, journalists, and the general public can examine and expose ads. This is a noble idea and one that deserves exploration. But the very fact that campaigns can make hundreds of versions of the same ad and that so much political propaganda comes from people and firms outside the traditional campaign and party system would render such a repository limited in its effects. In the absence of a legal requirement, companies have no reason to agree to and coordinate with such an archive. Beyond that, the U.S. Congress is unlikely to reform a system that their election campaigns are just learning to master.⁷

Facebook has pledged to install better filtering systems using artificial intelligence and machine learning to flag accounts that are run by bots, fake accounts (run by people who misrepresent their identities or interests), or accounts that otherwise violate Facebook's terms of service. These are just new versions of the technologies that have caused the problem in

the first place. And there would be no accountability beyond Facebook's word. The fact remains that in the arms race to keep propaganda flowing, human beings only review troublesome accounts long after the damage has been done. The prospect is stronger for reform in Europe and the United Kingdom. In 2017 the Information Commissioner's Office in the United Kingdom launched an investigation into the role played by Facebook and its use of citizens' data in the 2016 Brexit referendum and 2017 national elections.⁸

DISINFORMATION

We are in the midst of a worldwide, internet-based assault on democracy. Scholars at the Oxford Internet Institute have tracked armies of volunteers and "bots," or automated profiles, as they move propaganda across Facebook and Twitter in efforts to undermine trust in democracy or to elect favored candidates in the Philippines, India, France, the Netherlands, the United Kingdom, and elsewhere. We now know that agents in Russia are exploiting the powerful Facebook advertising system directly. Russian disinformation has worked its way into the social media feeds of voters in France, the Netherlands, and the United Kingdom in the past two years. Armies of volunteer social media activists, working on behalf of authoritarian-minded parties around the world, have flooded Facebook, Twitter, and Facebook-owned Instagram and WhatsApp with disinformation, propaganda, and threats against critics and journalists.⁹

This disinformation takes many forms and has many different motivations. Some of it was designed to attract clicks to generate advertising revenue. Some of it was designed by a political party, government, or nonstate actor to generate political pressure, to undermine trust in the institutions that make democracy function, or to disrupt democratic deliberation. And some of it seems to have been created just for the entertainment value of those who promote it. It's a mess, and it's difficult to discuss the phenomena as a single field or subject. Investigations of the various modes of disinformation require a wide array of scholarly and journalistic tools, including ethnography, data science, and digital forensics. Facebook itself, with its

global span and rich collections of user and advertising data, is the one actor that could help us accurately assess the impact of these behaviors on publics around the world. But Facebook has the most to lose if the data reveal that the company could have stemmed the flow of disinformation long ago and did nothing, or that—more likely—there is nothing Facebook can do to fix the problem, that it's endemic to the very core design of Facebook. So Facebook remains silent, leaving the rest of us to fumble around, trying to make sense of what seem to be rather severe threats.¹⁰

Facebook is hardly the only domain of disinformation. And it's rarely the birthplace of it. Other platforms such as 4Chan and 8Chan have been known to launch many of the most troubling campaigns of harassment and disinformation. Reddit, blogs, and websites that serve far-right nationalist groups often sprout it as well. Profiteering websites in places such as Macedonia have produced thousands of bogus stories and used Facebook and Twitter to attract readers to their sites, where advertising services generated revenue for them. Twitter is infected with bots that amplify false narratives and other forms of disinformation or distract critics of authoritarian leaders by harassing them with garbage.

To grasp how disinformation works in the world we must consider the entire ecosystem. Those who push this content start from the fringes, such as Reddit and 4Chan, where publication is easy and they can test various forms of their message among like-minded and motivated peers. They can generate catchy new hashtags for Twitter or join new content with established and popular hashtags.

They can post videos on YouTube, which then allows for easy redistribution via Twitter and Facebook, and which, thanks to Google's search algorithms favoring its own services, is a powerful social distribution platform in itself. Once items or issues rise in visibility on Twitter, editors of established online news and commentary sites such as BuzzFeed, Breitbart, Salon, or HuffPost can take notice. These sites are often starving for lively content that can connect to the hot issue or controversy of the day. So editors eagerly push young, underpaid reporters to echo each other with "hot takes" or sly commentary on whatever that issue might be—even to deflate or debunk the claim. It hardly matters if the story is important or even true

once this level of content-hungry news site amplifies the story even more. Each of these rather recently established news organizations has optimized its headline writing, image placement, and writing style for social media distribution. If they have not noticed the items already, more established news services such as the *Guardian*, the BBC, Fox News, CNN, and the *Washington Post* soon catch on to the buzz of the day and might prepare new versions of the story as well—again, often to debunk claims. But it hardly matters. At this point, the specious content has moved completely up the chain of the media ecosystem. At every link in the chain, Facebook, the largest and most powerful media system in the world, plays a role in amplifying the disinformation. But by the time the content reaches the level of HuffPost or Breitbart, it's irresistible to many Facebook users. Those who buy into the claims of the disinformation share it gleefully. Those who are appalled by the existence of the disinformation share it out of disgust. Whether intended to signal acclimation or disgust, comments and sharing work the same way and generate the same result. Facebook reads both negative and positive comments as "meaningful engagement," so it amplifies the message, pushing it into more News Feeds and making it appear higher up in the feed and more frequently. The effect is the same: chaos reigns. And the disinformation artists laugh at just how easy it all is.¹¹

The media ecosystem includes humans and machines, minds and algorithms, and operates through a series of similarly distinct and innocent choices: read this; report on this; share this; comment on this; click on this. As danah boyd, president of Data and Society, put it, "We've built an information ecosystem where information can fly through social networks (both technical and personal)." Having followed disinformation and harassment campaigns that emerge from gaming, misogynist cabals, white nationalist groups, and pornography internet subcultures for years, boyd warned that the technological interventions that Facebook promised to introduce would do little to stem the flow. Facebook amplifies and enables the phenomenon. But it originates among us—particularly among the most alienated of us. We built an attention economy, boyd argues. So we should expect people and groups to exploit that economy—ruthlessly and mercilessly.¹²

Beyond the algorithm, the very fact that we choose to perform our identities for others on Facebook amplifies the potential for disinformation to spread. “Sharing” an item on Facebook is an act that is both social and performative. Sharing is a declaration of identity. Sharing items about the San Antonio Spurs or FC Barcelona places me in particular circles. It also separates me from indignant fans of rival sports clubs. The act of sharing divisive material bonds me with some and separates me from others. But mostly it defines me to my circle of Friends.

So while Mark Zuckerberg thought he was forging a social network to connect people, by encouraging us to share content from other sites so easily he actually divided us by connecting us. These divisions are fluid, and some matter more than others. But over time, as Facebook structures our feeds to reward those who interact most frequently with us, our tribes solidify. Because we yearn for those small bolts of affirmation—the comment, the like, the share—we habitually post items that have generated the most response. We teach ourselves what sort of material will satisfy our tribe and will generate applause. Facebook also rewards us for that and pushes that rewarding content out farther, faster, and more frequently. If the item is false, hateful, or completely absurd, it hardly matters to the community. In fact, highly disputable, divisive, or disreputable content can become even more valuable as a signal of identity. Even if one loses Friends over it, or perhaps especially if one loses Friends over it, edgy or false or hateful content has the power of certifying that the person who posted it cares more for the identity marker than for the relationship.¹³

quickly—and unfortunately—dubbed the phenomenon “fake news.” BuzzFeed and its lead reporter Craig Silverman generated many of the most notable accounts of the “fake news” phenomenon. Silverman had been tracking the rise of such sites since at least 2014.

One of the keys to the success of “fake news” is that often these pieces were designed expertly to play both to the established habits of rapid sharers of Facebook content and to Facebook’s EdgeRank algorithm. They reinforce existing beliefs among a highly motivated subset of Facebook users. Absurd or controversial posts are likely to be shared and cheered by those willing to believe them and dismissed, commented upon, argued about, and shared by those who dismiss the veracity of those posts. If someone sees an obviously fraudulent claim on a Friend’s Facebook site and responds to it, it’s likely to flare a long and angry argument among different camps. As we know all too well, Facebook is designed to amplify that sort of engagement. So the pieces spread. BuzzFeed was the ideal news organization to dig deeply into this dynamic, as it was founded to generate similar levels of engagement for both its news posts and its lifestyle features, listicles, and quizzes.¹⁴

However, the right-wing mediasphere quickly tried to discredit Silverman, BuzzFeed, and the term “fake news.” An article in the conservative magazine *National Review* questioned Silverman’s methodology when he tried to determine if “fake news” stories had been shared more heavily than more traditional sources of news. And the ultranationalist site Breitbart ran a story accusing Silverman of manufacturing a moral panic around “fake news” to advance his career.¹⁵

Despite these efforts to undermine the reporting on the proliferation, the epithet of “fake news” carried some meaning for about six months, until about January 2017, when the Trump transition team and his supporters co-opted the term by using it repeatedly to characterize news produced by professional organizations that operated within the traditions of verification and correction. In other words, Trump and his band flipped the meaning of “fake news” almost completely, making it hard for any serious examination of the phenomenon to start with that term. In early January 2017 *Washington Post* columnist Margaret Sullivan, formerly the public editor of

THE DISINFORMATION ABOUT “FAKE NEWS”

As the 2016 U.S. election campaign moved into the fall, American journalists became aware of a proliferation of fiction circulating on Facebook, disguised as legitimate news. Many, but not all, of the items spread false stories that could have helped expand or solidify support for Donald Trump, including false claims about Hillary Clinton, Muslims, or Mexican immigrants. A notorious case involved a frequently shared item claiming Pope Francis had endorsed Trump for president. American journalists

the *New York Times* and editor-in-chief of the *Buffalo News*, declared the term useless and meaningless. “Faster than you could say ‘Pizzagate,’ the label has been co-opted to mean any number of completely different things: Liberal claptrap. Or opinion from left-of-center. Or simply anything in the realm of news that the observer doesn’t like to hear,” Sullivan wrote.¹⁶

“Pizzagate” was a term applied to a weird story that circulated among extreme right-wing news sites and social media users in the United States. It referred to a rumor that the email that had been stolen from the computer of John Podesta, a top Hillary Clinton campaign official, contained records of a child-sex ring run out of the basement of a pizza restaurant in Washington, D.C. Believing the false stories, a man bearing a gun entered the pizza restaurant in December 2016 to investigate and free the abused children he was convinced were being held in the nonexistent basement.¹⁷

The term “fake news” never covered the problem adequately in the first place. Not all of the troublesome items on Facebook were purely false. Sometimes they included elements of truth, even links to respectable news outlets, but took those elements and framed them in distorted or inaccurate ways to promote a partisan agenda. This is a classic technique of propaganda. In that way, this phenomenon was not new. Only its alarming amplification by its most fertile medium was new.

If the term “fake news” does not do the job, because it has been co-opted by the very forces that seek to undermine legitimate, professional news reporting practices and institutions, what terms can capture and describe what is going on? Searching for a better way to describe the variety of troublesome reports and doctored images floating down News Feeds, Caroline Jack, a researcher at the New York-based think tank Data and Society, set out to categorize and clarify the problem. Her report, entitled, “Lexicon of Lies: Terms for Problematic Information,” walks through the pros and cons, strengths and limitations of terms such as “propaganda,” “agitprop,” and “misinformation.” Jack described “disinformation” as “information that is deliberately false or misleading.” Disinformation has the added feature of being often irresistible to the clicking fingers of Facebook users and thus the distributive power of Facebook algorithms. After all, the most incredible stories are the most emotionally powerful despite

being—literally—incredible. The motives for such misleading can be pranksterish, selfish, or just malevolent. It seems that “disinformation” describes the widest set of phenomena with the most accurate description.¹⁸

The obsession with identifying, filtering, and quashing “fake news” distract many from the larger, deeper problem of disinformation. “Fake news” is a trivial problem compared with the general assaults on civic and democratic norms that have been at work for many years in the United States and around the world. Whether for laughs, blood, or profit, disinformation campaigns share a goal: to undermine trust in civic norms and institutions. Disinformation divides and debilitates a polis. With enough exposure over enough time, the idea of trust becomes comical, the idea of truth becomes irrelevant, and the idea of justice melds with tribal vengeance or retribution. Society devolves into a craven, selfish state in which, as in the title of journalist Peter Pomerantsev’s revealing book about Russia under Putin, “nothing is true and everything is possible.”¹⁹

THE AUTHORITARIAN PLAYBOOK

If you wanted to design a media system to support authoritarian leaders and antidemocratic movements, you could not do much better than Facebook. Katy Pearce, a scholar who studies social media use in Azerbaijan and other post-Soviet states, has deflated the idea that Facebook chiefly serves the cause of liberation and the promotion of democracy. From the 1980s right through the years during which Secretary of State Hillary Clinton promoted her “internet freedom” agenda, pundits and politicians promised that the introduction of digital media would create cracks in the systems of information control that authoritarian leaders had traditionally used to maintain power. The theory held that by connecting dissidents and nascent civil society movement (churches, labor unions, human rights groups) to allies and information from democratic societies, the former group would grow larger, bolder, and more effective. Even Ronald Reagan, who served as president a decade before the internet achieved global recognition, said, “The Goliath of totalitarianism will be brought down by the David of the microchip.” This theory has not held up to empirical scrutiny,

Pearce argues, Social media, especially Facebook, are certainly designed to aid civil society to organize efficiently, but they are also designed for authoritarian regimes to exploit with better resources than their opponents have.²⁰ There are five major ways that authoritarian regimes exploit Facebook and other social media services. They use it to organize countermovements to emerging civil society or protest movements. They can frame the public debate along their terms by virtue of having greater resources and technical expertise than their opponents and critics. They can allow or even manage social media as a stage for citizens to voice complaints without direct appeal or protest, thus letting citizens vent about corruption or governmental incompetence (as China allows citizens to use WeChat). Regimes can use social media to coordinate among elites to rally support as well.²¹

The fifth and most pernicious way authoritarian regimes use social media—especially Facebook—is in the surveillance and harassment of opposition activists and journalists. It's easy to plant a fake "Friend" with an activist group or an activist's profile page, allowing access to rich personal and organizational information. It's also easy to doctor photographs and video purporting to show an activist or journalist in a compromising situation, thus discrediting or at least distracting her or him and undermining the reformist effort. Once a sensational item about a well-known person arrives on Facebook, it gets around. And again, debunking it only enhances the story's penetration and influence. While most of the world was focusing on how social media helped insurgent groups in what became known as the 2011 "Arab Spring," the governments of Bahrain and Syria quashed nascent protest movements and thus avoided being grouped in with Tunisia and Egypt by deploying these tactics.²²

Vladimir Putin's regime has mastered these tactics and deployed them both domestically and internationally. Much of the pro-Trump and anti-immigrant material that showed up on American Facebook News Feeds came from a St. Petersburg-based company called the Internet Research Agency, which employs hundreds of people to generate and spread disinformation that could serve the interests of the Russian government. And the Russian state propaganda outlets RT and Sputnik have positioned themselves to seed anti-immigrant, anti-Muslim, and anti-establishment

disinformation into Germany, the United Kingdom, France, Ukraine, and the United States in a concerted effort to undermine trust in journalistic, governmental, and civil society institutions. Domestically, Putin's allies have spread abusive and harassing messages about critics, dissidents, and journalists. The Russian success at these efforts has made it a model for other authoritarian leaders and has brought attention to the matter within the United States and Western Europe. But the real masters and early adopters of pro-authoritarian social media use now rule two more populous countries than Russia: India and the Philippines.²³

The rise and consolidation of the Bharatiya Janata Party, or BJP, under Prime Minister Modi has been historic. Earlier in its history the BJP had formed weak coalition governments in India, but never before did it have the overwhelming parliamentary majority that it gained in the 2014 elections. The party had solidified a coalition of regional ethnic and religious nationalist parties that it called the National Democratic Alliance. No party had controlled a clear majority in Parliament since 1984, when India was essentially a one-party state under the control of the Congress Party of Indira Gandhi and her sons. The overwhelming BJP victory surprised pollsters and pundits. Much credit for it went to reports of endemic corruption within the Congress Party. But Modi's fervent embrace of social media did not go uncredited.²⁴

The BJP originated as the political wing of a militant Hindu nationalist movement called Rashtriya Swayamsevak Sangh (RSS). The RSS has long insisted on traditional roles for women, restricting immigration from majority-Muslim Bangladesh, a ban on celebrations of Valentine's Day because it signifies a degradation of morals by promoting "Western" ideas of love, and prohibitions on the slaughter and consumption of beef. One of its most fervent followers assassinated Mohandas Gandhi in 1948 because Gandhi and the Congress Party embraced a secular vision for India and tolerance of the 20 percent of the country that was Muslim. The RSS is the chief promoter of the idea of Hindutva, a protofascist vision of India as a Hindu homeland and theocracy. It has promoted a cleansing of history books, harassment of historians who describe India's multifaith past and traditions of tolerance, and a call for Hindu men to be strong and virulent in the defense

of their identity. Anti-Islamic politics and policies have been central to the RSS and thus the BJP. When Modi, a longtime member of the RSS, was chief minister of the state of Gujarat in 2002, an anti-Muslim pogrom broke out. More than 1,000 people died, 790 of who were Muslim. Although cleared of malfeasance by a subsequent investigation, Modi clearly failed to act on behalf of his Muslim citizens, and other BJP officials participated in the pogrom. Human Rights Watch accused Modi's government of covering up the government's role in the atrocities. The United States and the United Kingdom banned Modi from entering their countries until he assumed the prime ministership.²⁵

For a candidate so historically embedded within violent, nationalist forces, Facebook was the ideal platform to motivate potential supporters. The most caustic messages could fly below the view of journalists and international observers. Messages could rile up anti-Muslim passions and channel people to the polls. Not only did Modi develop a formidable social media team for his campaign, but his BJP continued to run a social media team staffed by both professionals and volunteers after the 2014 victory. The chief purpose of this team was to spread propaganda in favor of BJP policies. But it also had orders to destroy the reputations of journalists, civil society activists, critics of anti-Islam policies, and political enemies. In her book *I Am a Troll*, journalist Swati Chaturvedi tells the story of a woman who became enchanted with the BJP's Hindu nationalism while living in the United States. Upon her return to India, the woman enrolled as a social media worker at the BJP digital headquarters. At first she gleefully spread pro-BJP and anti-Congress Party items via WhatsApp. But after a while the negative messages struck her as cruel and unwarranted. "It was a never ending drip feed of hate and bigotry against minorities, the Gandhi family, the journalists on the hit list, liberals... anyone perceived as anti-Modi," the woman told Chaturvedi. "I simply could not follow [the social media director's] directions any more when I saw rape threats being made against female journalists like Barkha Dutt." The woman quit the BJP soon after.²⁶

Facebook staff worked with BJP officials during the Modi campaign. India has more Facebook users than any other country, with more than

250 million in 2018, about 30 million more than the United States has. That 250 million is less than one-quarter of the population of India, while the 220 million Americans on the service constitute more than 60 percent of the U.S. population. So not only is the future of Facebook in India, the present of Facebook is as well. Modi's Facebook page has more followers than that of any other world leader, 43 million, or almost twice what Donald Trump's Facebook page has.²⁷

Rival parties in India have established similar social media teams to mimic the BJP success. And now a slew of independent "troll farms" offer their services to private citizens, politicians, and companies that wish to destroy people's reputations. These services create fake videos of targets engaging in sex, drug use, or religious desecration. As the BJP solidifies its political standing the overall political culture of India has degraded through social media harassment and intimidation.²⁸

FREE BASICS ENABLES OPPRESSION

While Narendra Modi has managed to distance his personal reputation from the damage done by his supporters, the president of the Philippines, Rodrigo Duterte, has ridden his army of social media warriors from the position of longtime mayor of Davao City to president in May 2016. At each step, Duterte and his supporters engaged in virulent character assassination, threats, and harassment. Unlike India, the Philippines in 2015 invited Facebook to spread its Free Basics service. That means for the half of its 103 million citizens who regularly partake of social media on their mobile phones, Facebook use does not count against the monthly data allotment they pay for. Visiting reputable news sites costs data and thus money.²⁹

"We're one step closer to connecting the world as we launched Internet.org in the Philippines today," Zuckerberg wrote in March 2015. "Now everyone in the country can have free access to internet services for health, education, jobs and communication on the Smart network. Here's a photo of Jaime, a driver in Manila who uses Facebook and the internet to stay in touch with loved ones who moved to Dubai." It was a sweet, lovely post. Facebook was bringing Jaime's family closer together, at least virtually, and

on Facebook's terms. Zuckerberg showed no awareness that by unleashing this service on the Philippines he was inviting brutality and misery to descend on innocent Filipinos.³⁰

The dominance of Facebook in the media diet of Filipinos because of Free Basics was instant and timed perfectly for Duterte. The first move Duterte made when launching his campaign in 2015 was to hire a social media director and team. That team quickly leveraged the new connectivity created by Facebook's Free Basics service. Duterte's team of paid social media supporters, up to 500 volunteers, and thousands of bots have manufactured and spread false stories, and undermine trust in professional journalists. They deploy fake accounts to multiply the effects of the disinformation they share.³¹

Duterte's social media army during the 2016 campaign was divided into four groups, and each group was given a group to target: overseas Filipino workers, residents of the island of Luzon, Visayas, and Mindanao. Each group created its own content to express the daily narrative dictated by the Duterte campaign headquarters. Often these teams targeted university students they identified as Duterte critics. They shared and released students' mobile phone numbers. One Facebook post threatened a student with death. Other students received rape threats. Once Duterte assumed office he proceeded to shun interviews and encounters with professional journalists, choosing instead to communicate via Twitter and Facebook.³²

Facebook itself made all of this possible. In January 2016 Facebook sent three employees to Manila to train the various presidential candidates and their staffs how best to use the service. The Facebook team met Duterte's campaign staff at the Peninsula Manila Hotel. The campaign staff learned the basics of setting up a campaign page, getting it authenticated with a blue check mark, and attracting followers. Duterte's campaign constructed a social media apparatus unlike that of any other candidate in the race. Every day the campaign would create a message for the next day. The teams of volunteers would then pump the message across both real and fake Facebook accounts, some with hundreds of thousands of followers. Facebook almost immediately received complaints about inauthentic pages. Soon, however, the company fielded complaints about Duterte's Facebook army

circulating insults and violent threats. The Duterte campaign instructed its followers to circulate bogus stories, including a false endorsement by Pope Francis, with the phrase "Even the Pope Admires Duterte" pasted under the pope's image. Truth ceased to matter. Following a pattern familiar to those who watch politics in India, Russia, Ukraine, and even Estonia, the troll-army-driven propaganda dominated the political conversation. Any discussion of policies or choices or compromises became impossible. Opponents, journalists, and civil society leaders were left gasping at the audacity of Duterte's stunts. By April 2016, just one month before the election, a Facebook report called him the "undisputed king of Facebook conversations." Duterte occupied 64 percent of all election-related conversations on the Facebook pages in the Philippines.³³

After Duterte won, Facebook extended its partnership with the administration, helping Duterte execute his violent, nationalist agenda. Duterte banned the independent press from covering his inauguration live from inside Rizal Ceremonial Hall. He didn't need journalists. He just had the inaugural events streamed live on Facebook. With the rise of Duterte, Facebook solidified itself as the only media service that matters in the Philippines, a country of more than 105 million people, rich with resources, a multilingual and globally dispersed population, and a rich history of anti-colonial resistance.³⁴

Duterte's direct communication with his supporters through social media has been powerful and effective. Since his election, Duterte has openly fostered a culture of vigilantism. More than fourteen hundred people have been killed by police and civilians because they were accused or suspected of selling or using illegal drugs. The Philippines was not long ago known as a country with a strong civil society and stable democratic values. A popular uprising in 1986 called People's Power expelled corrupt dictator Ferdinand Marcos and installed reformer Cory Aquino as president. People Power in the Philippines, following the fall of the military junta in Brazil in 1985, was among the first of a series of largely nonviolent revolutions between 1985 and 1991 that swept away authoritarian governments from South Africa to Poland to the Soviet Union. Filipinos rose up again in 2001 to peacefully reject a corrupt leader, President Joseph Estrada.

At the time, the 2001 uprising was among the first during which activists used SMS text messaging to coordinate the movement. Since 2001, the political culture of the Philippines has degraded. But it took its most serious downturn just as Facebook introduced Free Basics and Duterte took full advantage of it.³⁵

In November 2017 Facebook announced a new partnership with the Duterte regime. Facebook will work with the government to lay underwater data cables that will bypass the Luzon Strait, where typhoons and earthquakes have often damaged standard cables. Facebook will fund the underwater links. The government will build cable landing stations. The Philippines has been moving for some years to become the central hub of optical fiber cables to facilitate digital data flows for East Asia and the South Pacific. A deep and profitable partnership with Facebook, while Facebook serves as the chief propaganda and harassment platform for the Duterte regime, means that Facebook will not have much choice but to continue to support Duterte as he expands his campaigns of terror.³⁶

Like Duterte, Hun Sen, the dictator of Cambodia, has become a Facebook star as he has leveraged the power of Free Basics to harass his opponents and promote his image. He has used the classic authoritarian playbook: develop a following; ensure that independent media can't compete with state propaganda on Facebook; make sure Facebook is the equivalent of the internet itself; and employ a troll army (in Hun's case, hired from firms based in India and the Philippines—two countries with experience in just such methods) both to push items that show him in a positive light and to terrorize and humiliate opponents and critics. Most important, Hun's staff works directly with Facebook staff to silence critics and maximize the influence Hun's Facebook pages can generate.³⁷

In October 2017 Facebook changed how news provided by professional news services would appear on the pages of users in Cambodia, Slovakia, Sri Lanka, Bolivia, Guatemala, and Serbia. Professional and independent news items would no longer run in the main News Feed along with advertisements, personal posts, and music videos. Instead news would sit on a separate, harder-to-see tab on the Facebook page. The results were predictable. Traffic from Facebook to independent news sites

in all of these countries plummeted. Hun got an even better experience from Facebook. In angry response to this change, Serbian journalist Stevan Dojcinovic wrote an op-ed in the *New York Times* called, “Hey, Mark Zuckerberg: My Democracy Isn’t Your Laboratory.” In it, Dojcinovic wrote, “The major TV channels, mainstream newspapers and organized-crime-run outlets will have no trouble buying Facebook ads or finding other ways to reach their audiences. It’s small, alternative organizations like mine that will suffer.” In every country in which Facebook made this change, Facebook dominates the media ecosystem. Even small changes in Facebook’s design or algorithmic emphasis can alter the political fortunes of an entire nation. Zuckerberg presented Internet.org and Free Basics as tools that would bring people together. The malevolent consequences of this benevolent concept have been all too clear. Facebook canceled the journalism experiment after the uproar. But its willingness to use small, poor countries as laboratories should trouble anyone who cares about global peace and stability.³⁸

In perhaps the most alarming account of how Facebook has further empowered violent, oppressive state power, Free Basics has also transformed Myanmar, a country that has just recently emerged from decades of military rule. In 2014, only a year before the country’s first free elections since 1960 elevated the human rights advocate Aung San Suu Kyi to power, the first mobile phone companies established data connections in Myanmar. Government policy pushed for rapid adoption of mobile phones. And by 2016 Facebook introduced Free Basics, ensuring that—like in the Philippines—for many Burmese there was no distinction between Facebook and the internet. To be connected means to use Facebook. Any other service costs money. The country has not had time to develop a mature and professional media system or a tradition of professional journalism. So the chaotic nature of Facebook is even more pronounced than in most of the rest of the world. In a country so recently opened up after more than fifty years of military control of every facet of daily life, Facebook seems like a gift. Now Burmese people can read and share news and rumors, experience music videos, share jokes, and, of course, spread hate.³⁹ Under military rule, rumor was the dominant form of “news” and the chief subject of discussion. So through Facebook, old habits thrive.

Buddhist nationalists have spread rumors of a global Muslim conspiracy bent on ridding the world of Buddhism. Through Facebook they have called for boycotts of Muslim-owned businesses, a ban on interfaith marriages, and limitations on rights for Muslims who live in Myanmar. Anti-Muslim riots broke out in 2015 in cities across Myanmar. By 2017 Buddhist attackers, supported by the military, carried out genocidal attacks on the Muslim Rohingya minority in western Myanmar, driving many into neighboring Bangladesh, which had no ability to absorb them. Many Rohingya refugees proceeded west into India. In September 2017, as the world took notice of the genocide, Aung San Suu Kyi posted a message on her own Facebook page declaring the genocide a myth and blaming “terrorists” for unrest in western Myanmar. She then went to India to meet with Modi to enlist his help. Modi concurred with her assessment of the Muslim threat.⁴⁰

Atrocities such as the slaughter and expulsion of the Muslim Rohingya people of Myanmar are hardly new. They have been a constant in human history. And, like with the various uprisings of 2011, people are going to use the communicative technologies available to them. So we should not be surprised that authoritarians, religious bigots, and ethnic nationalists push out disinformation via Facebook. We must, however, carefully consider what specific features and structures of Facebook make it so useful in such a consistent way by authoritarians and the movements that support them. Facebook allows authoritarian leaders and nationalist movements to whip up sentiment and organize violence and harassment against enemies real and imagined. It’s like nothing before. Its ubiquity and ease of use in countries that are still struggling after centuries of colonial rule—Kenya, Philippines, Cambodia, and Myanmar—offer the most destructive forces an ideal propaganda system. Facebook does not favor hatred. But hatred favors Facebook.